

ML hw7

0756607 范姜良瑜

Principle component analysis (PCA)

Assume we have data of $N \times \text{dim}$. N is the number of samples and dim is the number of features for each sample. `mnist_X.csv` provides data of shape 5000×784 .

1. Center each feature data to zero-mean

```
x = X - (np.mean(X, axis=0))
```

2. Compute covariance matrix Σ to get eigenvalues and eigenvectors

$$\Sigma = \frac{1}{n} \sum_{i=1}^n (X_i^T X_i)$$

```
1 covariance_mat = np.cov(x.T)
2 eig_vals, eig_vecs = LA.eig(covariance_mat)
```

3. Select principle components

The eigenvectors with higher eigenvalues are more informative than those with low eigenvalues. As a result, we rank the eigenvectors with eigenvalues and select the top k eigenvectors.

```
1 eig_pairs = [(np.abs(eig_vals[i]), eig_vecs[:,i]) for i in range(len(eig_vals))]
2 eig_pairs = sorted(eig_pairs, key=lambda k: k[0], reverse=True)
```

We use **explained variance** to see how informative of each eigenvector (principle component).

```
1 eigv_sum = sum(eig_vals)
2 for i, j in enumerate(eig_pairs):
3     print('eigenvalue {0:}: {1:.2%}'.format(i+1, (j[0]/eigv_sum).real))
```

Sample output:

```
1 eigenvalue 1: 14.78%
2 eigenvalue 2: 8.17%
3 eigenvalue 3: 6.31%
4 eigenvalue 4: 5.99%
5 eigenvalue 5: 5.11%
6 eigenvalue 6: 3.88%
7 eigenvalue 7: 3.38%
8 eigenvalue 8: 2.66%
9 eigenvalue 9: 2.55%
10 eigenvalue 10: 2.19%
11 ...
```


Here, we project the data to 2D space, i.e. $k = 2$.

```
1 N, dim = x.shape
```

```

2 w = np.hstack((eig_pairs[0][1].reshape(dim,1), eig_pairs[1][1].reshape(dim,1)))
3 x_PCA = x.dot(w)
4 # plot
5 for l,c in zip(labels,('C0', 'C1', 'C2', 'C3', 'C4')):
6     target_data = x_PCA[y == l]
7     plt.scatter(target_data[:,0].real, target_data[:,1].real, c=c, s=7, alpha=0.5)
8 plt.legend(labels.astype(int), loc='upper left')
9 plt.show()

```




project 784-dimensional data to 2D space with PCA

Also, we try to map the data to 3D space by taking top 3 eigenvectors.

```

1 num_dims = 3
2 w = np.hstack((eig_pairs[0][1].reshape(dim,1), eig_pairs[1][1].reshape(dim,1)))
3 for i in range(2,num_dims):
4     w_3 = np.hstack((w, eig_pairs[i][1].reshape(dim,1)))
5 x_PCA_3 = x.dot(w_3)

```



project 784-dimensional data to 3D space with PCA

Linear discriminant analysis (LDA)

Assume we have data of $N \times \text{dim}$. N is the number of samples and dim is the number of features for each sample. `mnist_X.csv` provides data of shape 5000×784 .

1. Compute dim -dimensional mean vectors m_i

$$m_i = \frac{1}{n_i} \sum_{x \in D_i}^n x_k$$

```

1 mean_vectors = []
2 for l in labels:
3     # column mean, i.e. mean of each feature, in the same class
4     mean_vectors.append(np.mean(x[y == l], axis=0))

```

2. Compute scatter matrices

First, we compute the within-class scatter matrix S_w . S_i is the scatter matrix for each class.

$$S_w = \sum_{i=1}^c S_i$$

$$S_i = \sum_{x \in D_i}^n (x - m_i)(x - m_i)^T$$

```

1 dim = x.shape[1]
2 s_w = np.zeros((dim, dim))
3 for l, mv in zip(labels, mean_vectors):
4     scatter_mat = np.zeros((dim, dim))
5     for d in x[y == l]:
6         d = d.reshape(dim, 1)
7         mv = mv.reshape(dim, 1)
8         scatter_mat += (d - mv).dot((d - mv).T)
9     s_w += scatter_mat

```

Then, we compute the between-class scatter matrix S_b . N_i is the number of data with label i . m_i is the i^{th} mean in the mean vector. m is the overall means of each feature across classes.

$$S_b = \sum_{i=1}^c N_i(m_i - m)(m_i - m)^T$$

```

1 overall_mean = np.mean(x, axis=0)
2 s_b = np.zeros((dim, dim))
3 for l, mv in zip(labels, mean_vectors):
4     # n is the number of data with label l
5     n = x[y == l].shape[0]
6     overall_mean = overall_mean.reshape(dim, 1)
7     mv = mv.reshape(dim, 1)
8     s_b += n * (mv - overall_mean).dot((mv - overall_mean).T)

```

3. Eigendecomposition of $S_W^{-1}S_B$

```

1 # check if s_w is a singular matrix
2 if LA.det(s_w) == 0:
3     eig_vals, eig_vecs = LA.eig(LA.pinv(s_w).dot(s_b)) # use pseudo inverse
4 else:
5     eig_vals, eig_vecs = LA.eig(LA.inv(s_w).dot(s_b))

```


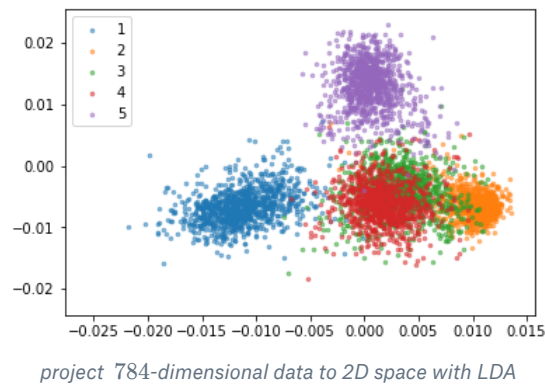
4. Select linear discriminants

Here, we select the top 2 linear discriminants to project the data onto a 2D space.

```

1 eig_pairs = [(np.abs(eig_vals[i]), eig_vecs[:,i]) for i in range(len(eig_vals))]
2 eig_pairs = sorted(eig_pairs, key=lambda k: k[0], reverse=True)
3 w = np.hstack((eig_pairs[0][1].reshape(dim,1), eig_pairs[1][1].reshape(dim,1)))
4 x_LDA = x.dot(w)
5 # plot
6 for l,c in zip(labels, ('C0', 'C1', 'C2', 'C3', 'C4')):
7     target_data = x_LDA[y == l]
8     plt.scatter(target_data[:,0].real, target_data[:,1].real, c=c, s=7, alpha=0.5)
9 plt.legend(labels.astype(int), loc='upper left')
10 plt.show()

```



Symmetric SNE and t-SNE

SNE

Convert the euclidean distances between data points to **conditional** probabilities that represent similarities.

$$p_{j|i} = \frac{\exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma_i^2}\right)}{\sum_{k \neq i} \exp\left(-\frac{\|x_i - x_k\|^2}{2\sigma_i^2}\right)}$$

The probability of point x_j being a neighbor of point x_i is proportional to the distance between these two points.

note: $p_{i|i} = 0$

y is an $N \times 2$ matrix that is our 2D representation of x .

$$q_{j|i} = \frac{\exp(-||y_i - y_j||^2)}{\sum_{k \neq i} \exp(-||y_i - y_k||^2)}$$

Goal: pick points in y such that q is similar to p .

Minimize the following cost by using gradient descent:

$$C = \sum_i KL(P_i || Q_i) = \sum_i \sum_j p_{j|i} \log \frac{p_{j|i}}{q_{j|i}}$$

Pick radius of Gaussian σ_i

The **perplexity** of any of the rows of the conditional probabilities matrix P is defined as:

$$Perp(P_i) = 2^{H(P_i)}$$

$$H(P_i) = - \sum_j p_{j|i} \log_2 p_{j|i}$$

Higher perplexity will increase the number of neighbors each point has.

Perform **binary search** over each σ_i until $Perp(P_i) =$ the desired perplexity.

Symmetric SNE

Minimize a KL divergence over the **joint** probability distributions with entries p_{ij} and q_{ij} .

$$p_{ij} = \frac{p_{i|j} + p_{j|i}}{2N}$$

$$q_{ij} = \frac{\exp(-||y_i - y_j||^2)}{\sum_{k \neq i} \exp(-||y_k - y_i||^2)}$$

```

1 def q_joint(Y):
2     """
3         Given low-dimensional representations Y, compute matrix of joint
4         probabilities with entries q_ij.
5     """
6     # compute distances from every y_i to y_j
7     dists = neg_squared_euclidean_dists(Y)
8     exp_dists = np.exp(dists)
9     # let q_ii = 0
10    np.fill_diagonal(exp_dists, 0.)
11    # divide by the sum of the exponential matrix
12    Q = exp_dists / np.sum(exp_dists)
13    return Q, None

```

Use the following gradient to update the i 'th row of our low-dimensional representation y .

$$\frac{\partial C}{\partial y_i} = 2 \sum_j (p_{ij} - q_{ij})(y_i - y_j)$$

note: the coefficient on the handout is 2, but some use 4 according to some online articles.

```

1 def grad_symmetric_sne(P, Q, Y, _):
2     pq_diff = P - Q                                # N x N
3     pq_expand = np.expand_dims(pq_diff, 2)          # N x N x 1
4     y_diff = np.expand_dims(Y, 1) - np.expand_dims(Y, 0) # N x N x 2
5     grad = 2. * (pq_expand * y_diff).sum(1)         # N x 2

```

```
6 return grad
```

t-SNE

In high dimension, Gaussian distribution is used to turn euclidean distances into probabilities. In **low dimension**, **Student's t-distribution** is used to alleviate crowding problem.

$$p_{ij} = \frac{\exp\left(-\|x_i - x_j\|^2 / 2\sigma_i^2\right)}{\sum_{k \neq l} \exp\left(-\|x_l - x_k\|^2 / 2\sigma_l^2\right)}$$

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}}$$

$$\frac{\partial C}{\partial y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j) (1 + \|y_i - y_j\|^2)^{-1}$$

Differences between symmetric SNE and t-SNE

In t-SNE, Student's t-distribution is to turn euclidean distances into probabilities in low dimension. As a result, the computation for gradients are also different.

symmetric SNE	t-SNE
$q_{ij} = \frac{\exp(-\ y_i - y_j\ ^2)}{\sum_{k \neq l} \exp(-\ y_k - y_l\ ^2)}$	$q_{ij} = \frac{(1 + \ y_i - y_j\ ^2)^{-1}}{\sum_{k \neq l} (1 + \ y_k - y_l\ ^2)^{-1}}$
$\frac{\partial C}{\partial y_i} = 2 \sum_j (p_{ij} - q_{ij})(y_i - y_j)$	$\frac{\partial C}{\partial y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j) (1 + \ y_i - y_j\ ^2)^{-1}$

Redefine the provided `tsne(.)` as `sne(.)`. Add new parameters `q_func` and `grad_func`. The default setting is symmetric SNE. The implementation of `q_joint(.)` and `grad_symmetric_sne(.)` is shown above.

```
1 def sne(X=np.array([]), no_dims=2, initial_dims=50, perplexity=30.0, q_func=q_joint, grad_func=grad_symmetric_sne):
2     """
3     Runs symmetric SNE or t-SNE on the dataset in the NxD array X to reduce its
4     dimensionality to no_dims dimensions. The syntaxis of the function is
5     `Y = tsne.tsne(X, no_dims, perplexity), where X is an NxD NumPy array.
6     """
7     ...
8     for iter in range(max_iter):
9         Q, dists = q_func(Y)
10        Q = np.maximum(Q, 1e-12)
11        # Compute gradient
12        dY = grad_func(P, Q, Y, dists)
13    ...
```


The following shows the iteration error of symmetric SNE.

```
1 Iteration 10: error is 26.346534
2 Iteration 20: error is 20.141809
3 Iteration 30: error is 17.839637
4 Iteration 40: error is 17.920184
5 Iteration 50: error is 17.982681
6 ...
7 Iteration 980: error is 2.288007
8 Iteration 990: error is 2.288007
```

```
9 Iteration 1000: error is 2.288007
```

The following shows the iteration error of t-SNE.

```
1 Iteration 10: error is 25.862092
2 Iteration 20: error is 22.084489
3 Iteration 30: error is 19.595014
4 Iteration 40: error is 18.567858
5 Iteration 50: error is 18.268595
6 ...
7 Iteration 980: error is 1.302603
8 Iteration 990: error is 1.301775
9 Iteration 1000: error is 1.300965
```



Eigenfaces

First, we print the first 50 images from `att_faces` to see how they look like.




the first 50 images from `att_faces`

Use PCA to find the top 25 principal components. Reshape the eigenvectors, i.e. principal components, to get the eigenfaces.

```
1 n_samples, h, w = images.shape
2 n_components = 25
3 centered_data, pc, mean = PCA(X, n_components)
4 eigenfaces = (pc.T).reshape((n_components, h, w))
5 plot_faces(eigenfaces.real, 5, 5)
```

```
1 def PCA(X, n_components):
2     # standardize data to ~ N(0,1)
3     N, dim = X.shape
4     # x = (x - x.mean()) / x.std()
5     mean = (np.mean(X, axis=0))
6     x = X - mean
7
8     # compute covariance matrix
9     covariance_mat = np.cov(x.T)
10    eig_vals, eig_vecs = LA.eig(covariance_mat)
11
12    eig_pairs = [(np.abs(eig_vals[i]), eig_vecs[:,i]) for i in range(len(eig_vals))]
13    eig_pairs = sorted(eig_pairs, key=lambda k: k[0], reverse=True)
14
15    w = np.hstack((eig_pairs[0][1].reshape(dim,1), eig_pairs[1][1].reshape(dim,1)))
16    for i in range(2, n_components):
17        w = np.hstack((w, eig_pairs[i][1].reshape(dim,1)))
18
19    return x, w, mean
```

```
1 def plot_faces(images, n_row, n_col):
2     plt.figure(figsize=(2.2 * n_col, 2.2 * n_row))
3     for i in range(n_row * n_col):
4         plt.subplot(n_row, n_col, i+1)
5         plt.imshow(images[i], cmap=plt.cm.gray)
6     plt.show()
```

Reconstruct the faces using the eigenfaces. Each face is a weighted combination of the eigenfaces. We dot the centered data and eigenfaces to get the weights. Weight the eigenfaces will give us the centered face and add it with the mean face to get the final reconstructed faces.

```

1 def reconstruct_face(centered_data, pc, mean, h, w, img_idx):
2     weights = np.dot(centered_data, pc.T)
3     centered_vector = np.dot(weights[img_idx, :], pc)
4     reconstructed_image = ( mean + centered_vector).reshape(h, w)
5
6     return reconstructed_image

```




reconstructed faces the first 50 images

Then, we randomly pick 10 images and reconstruct them.


```

1 img_idx = np.random.randint(images.shape[0], size=10)
2 img_samples = np.array([plt.imread(img) for img in (np.array(image_names)[img_idx])], dtype=
  p.float64)
3 n, n_col = 1, 10
4 plot_faces(img_samples, n, n_col)      # show the original images
5 recovered_images = [reconstruct_face(centered_data, pc.T.real, mean, h, w, i) for i in img_id
  x]
6 plot_faces(recovered_images, n, n_col) # show the reconstructed results

```



randomly sampled 10 original images



reconstructed the above randomly sampled images

Reference

PCA

- [機器/統計學習:主成分分析\(Principal Component Analysis, PCA\)](#)
- [Implementing a Principal Component Analysis \(PCA\) – in Python, step by step](#)

Eigenfaces

- [Eigenfaces: Recovering Humans from Ghosts](#)

LDA

- [Linear Discriminant Analysis – Bit by Bit](#)

SNE

- [Symmetric SNE and t-SNE](#)